

Accuracy vs compute, in-distribution test set (300 episodes, 95% bootstrap CI)

